

## CIS 530—Advanced Data Mining



# 8- Min Hashing and Locality Sensitive Hashing

Thomas W Gyeera, Assistant Professor  
Computer and Information Science  
University of Massachusetts Dartmouth

*Courtesy to Prof. Panayiotis Tsaparas*

# Why is similarity important?

---

- We saw many definitions of similarity and distance
- How do we make use of similarity in practice?
- What issues do we have to deal with?

# An important problem

---

- **Recommendation** systems
  - When a user buys an **item** (initially books) we want to recommend other items that the user may like
  - When a user rates a **movie**, we want to recommend movies that the user may like
  - When a user likes a **song**, we want to recommend other songs that they may like
- A big success of data mining
- Exploits the long tail

# Recommendation systems

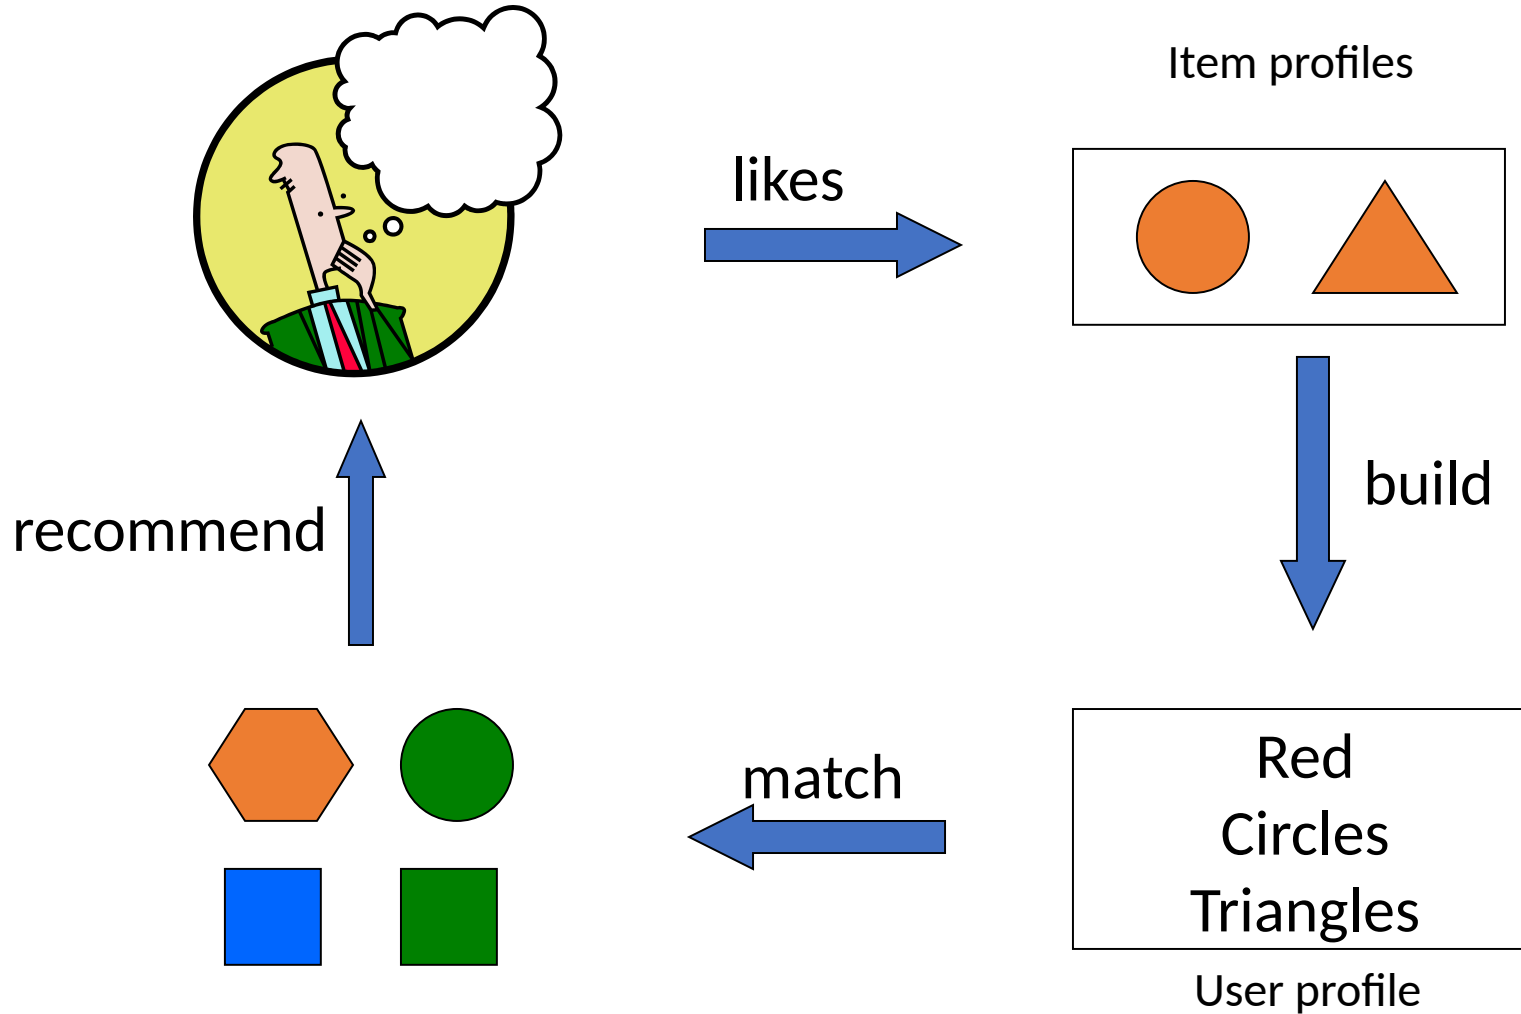
---

- **Content-based:**

- Represent the items into a **feature space** and recommend items to customer C **similar** to previous items rated highly by C
- Movie recommendations: recommend movies with same actor(s), director, genre, ...
- Websites, blogs, news: recommend other sites with “similar” content

# Plan of action

---



# Limitations of content-based approach

---

- Finding the appropriate features
  - e.g., images, movies, music
- Overspecialization
  - Never recommends items outside user's content profile
  - People might have multiple interests
- Recommendations for new users
  - How to build a profile?

# Recommendation Systems (II)

---

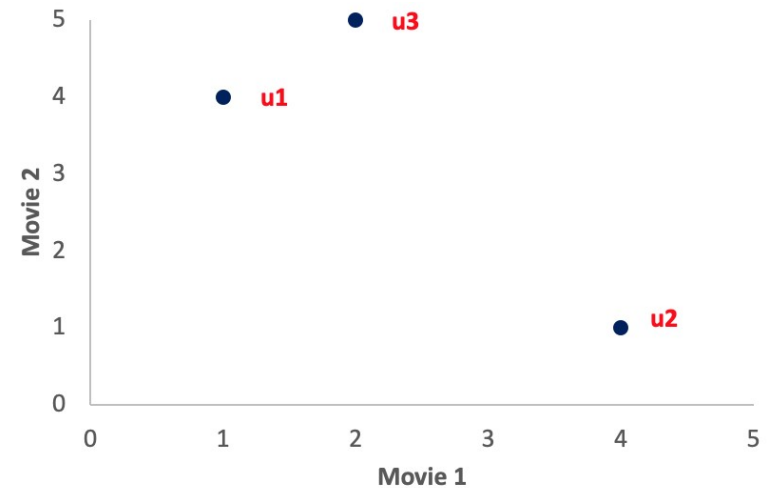
- Collaborative Filtering (user-user)
  - Consider user  $c$
  - Find set  $D$  of other users whose ratings are “similar” to  $c$ ’s ratings
  - Estimate user’s ratings based on ratings of users in  $D$

# Recommendation Systems (II)

---

- Collaborative Filtering (user-user)

	m1	m2
u1	1	4
u2	4	1
u3	2	5





# Recommendation Systems (III)

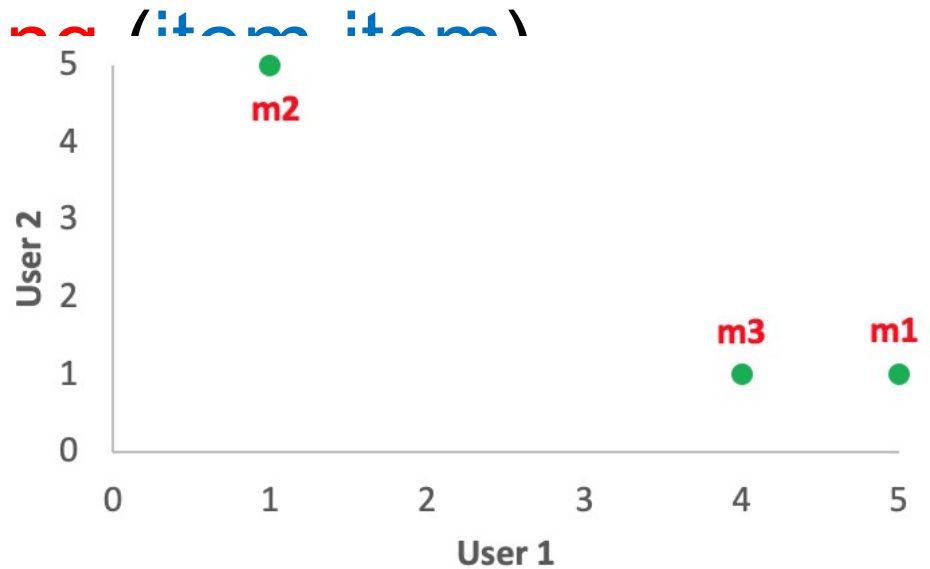
---

- Collaborative Filtering (item-item)
  - For item  $s$ , find other similar items
  - Estimate rating for item based on ratings for similar items
  - Can use same similarity metrics and prediction functions as in user-user model
- In practice, it has been observed that item-item often works better than user-user

# Recommendation Systems (III)

## • Collaborative Filtering

	u1	u2
m1	5	1
m2	1	5
m3	4	1



$$\cos(\theta) = \frac{A \cdot B}{|A||B|}$$



	m1	m2	m3
m1	1.000000	0.384615	0.998868
m2	0.384615	1.000000	0.428086
m3	0.998868	0.428086	1.000000

# Pros and cons of collaborative filtering

---

- Works for any kind of item
  - No feature selection needed
- New user problem
- New item problem
- Sparsity of rating matrix
  - Cluster-based smoothing?

# Another important problem

---

- Find **duplicate** and **near-duplicate** documents from a web crawl.
- Why is it important:
  - Identify **mirrored web pages**, and avoid indexing them, or serving them multiple times
  - Find **replicated news stories** and cluster them under a single story.
  - Identify plagiarism
- What if we wanted exact duplicates?

# Finding similar items

---

- Both the problems we described have a common component
  - We need a quick way to find **highly similar** items to a **query** item
  - OR, we need a method for finding **all pairs** of items that are **highly similar**.
- Also known as the **Nearest Neighbor** problem, or the **All Nearest Neighbors** problem
- We will examine it for the case of near-duplicate web documents.

# Main issues

---

- What is the **right representation** of the document when we check for similarity?
  - E.g., representing a document as a set of characters will not do (why?)
- When we have billions of documents, keeping the full text in memory is not an option.
  - We need to find a **shorter representation**
- How do we do **pairwise comparisons** of billions of documents?
  - If exact match was the issue, it would be ok, can we replicate this idea?

# Three Essential Techniques

---

1. **Shingling**: convert documents, emails, etc., to sets.
2. **Minhashing**: convert large sets to short signatures, while preserving similarity.
3. **Locality-Sensitive Hashing (LSH)**: focus on pairs of signatures likely to be similar.

# Motivating problem

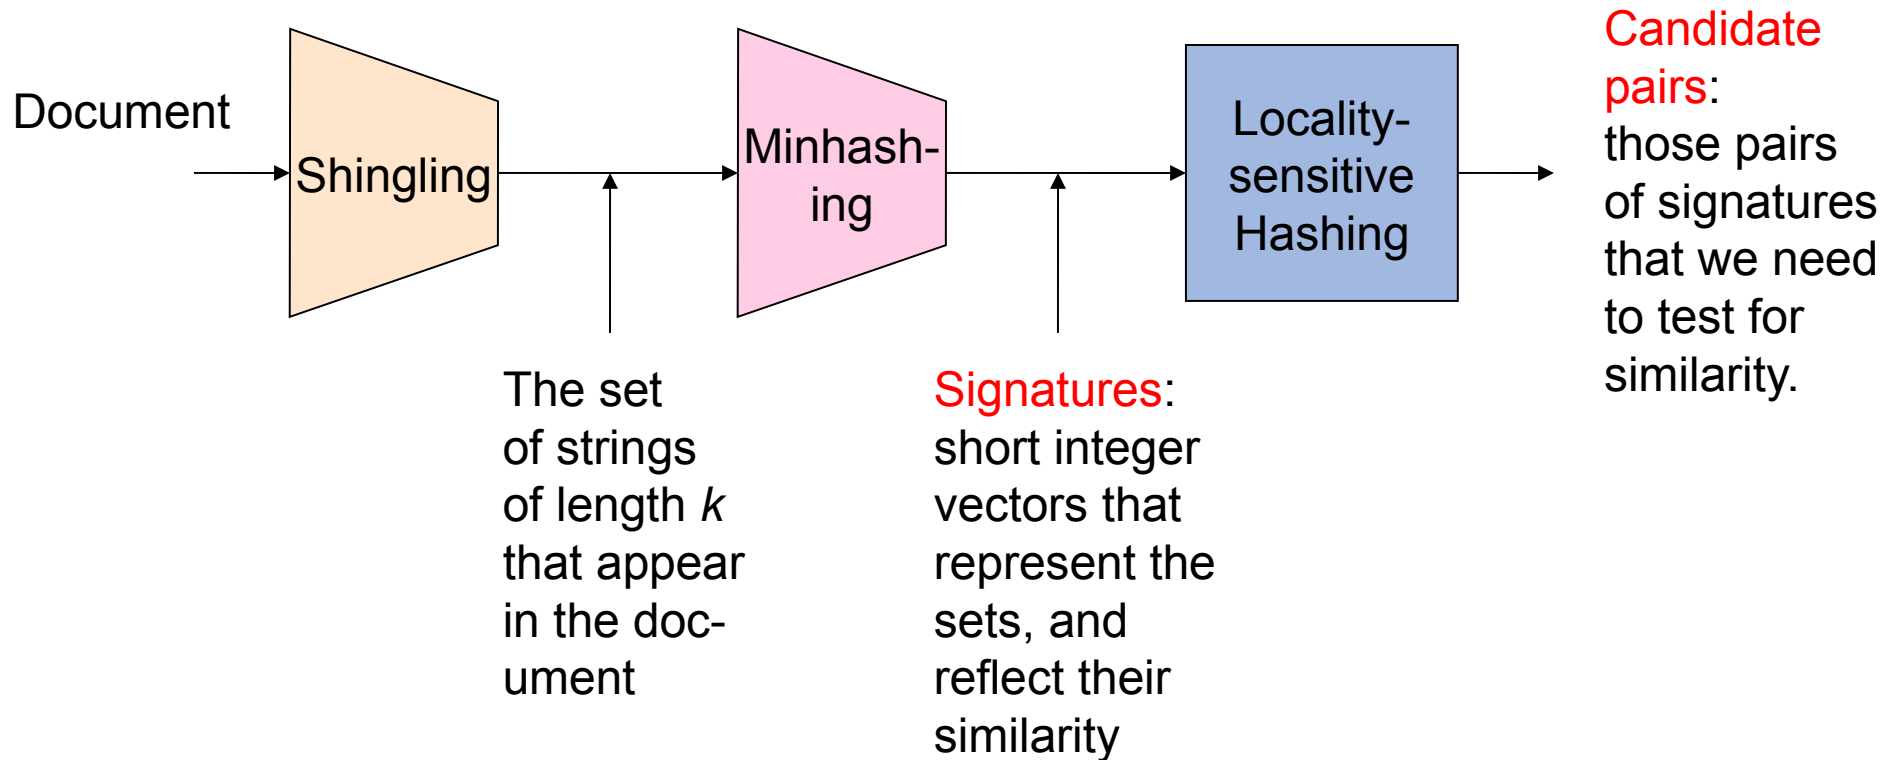
---

- Find **duplicate** and **near-duplicate** documents from a web crawl.
- If we wanted exact duplicates, we could do this by hashing
  - We will see how to adapt this technique for **near duplicate** documents



# The Big Picture

---



# Shingles

---

- A **k-shingle** (or **k-gram**) for a document is a sequence of **k** characters that appears in the document.
- **Example**: document = **ab****ca****b**. **k=2**
  - Set of 2-shingles = {**ab**, **bc**, **ca**}.
  - **Option**: regard shingles as a **bag**, and count **ab** twice.
- Represent a document by its set of **k**-shingles.

# Shingling

---

- Shingle: a sequence of  $k$  contiguous characters

a rose is a rose is a rose

a rose is

rose is a

rose is a

ose is a r

se is a ro

e is a ros

is a rose

is a rose

s a rose i

a rose is

a rose is

# Working Assumption

---

- Documents that have lots of shingles in common have similar text, even if the text appears in different order.
- **Careful**: you must pick  $k$  large enough, or most documents will have most shingles.
  - Extreme case  $k = 1$ : all documents are the same
  - $k = 5$  is OK for short documents;  $k = 10$  is better for long documents.
- Alternative ways to define shingles:
  - Use words instead of characters
  - Anchor on stop words (to avoid templates)

# Shingles: Compression Option

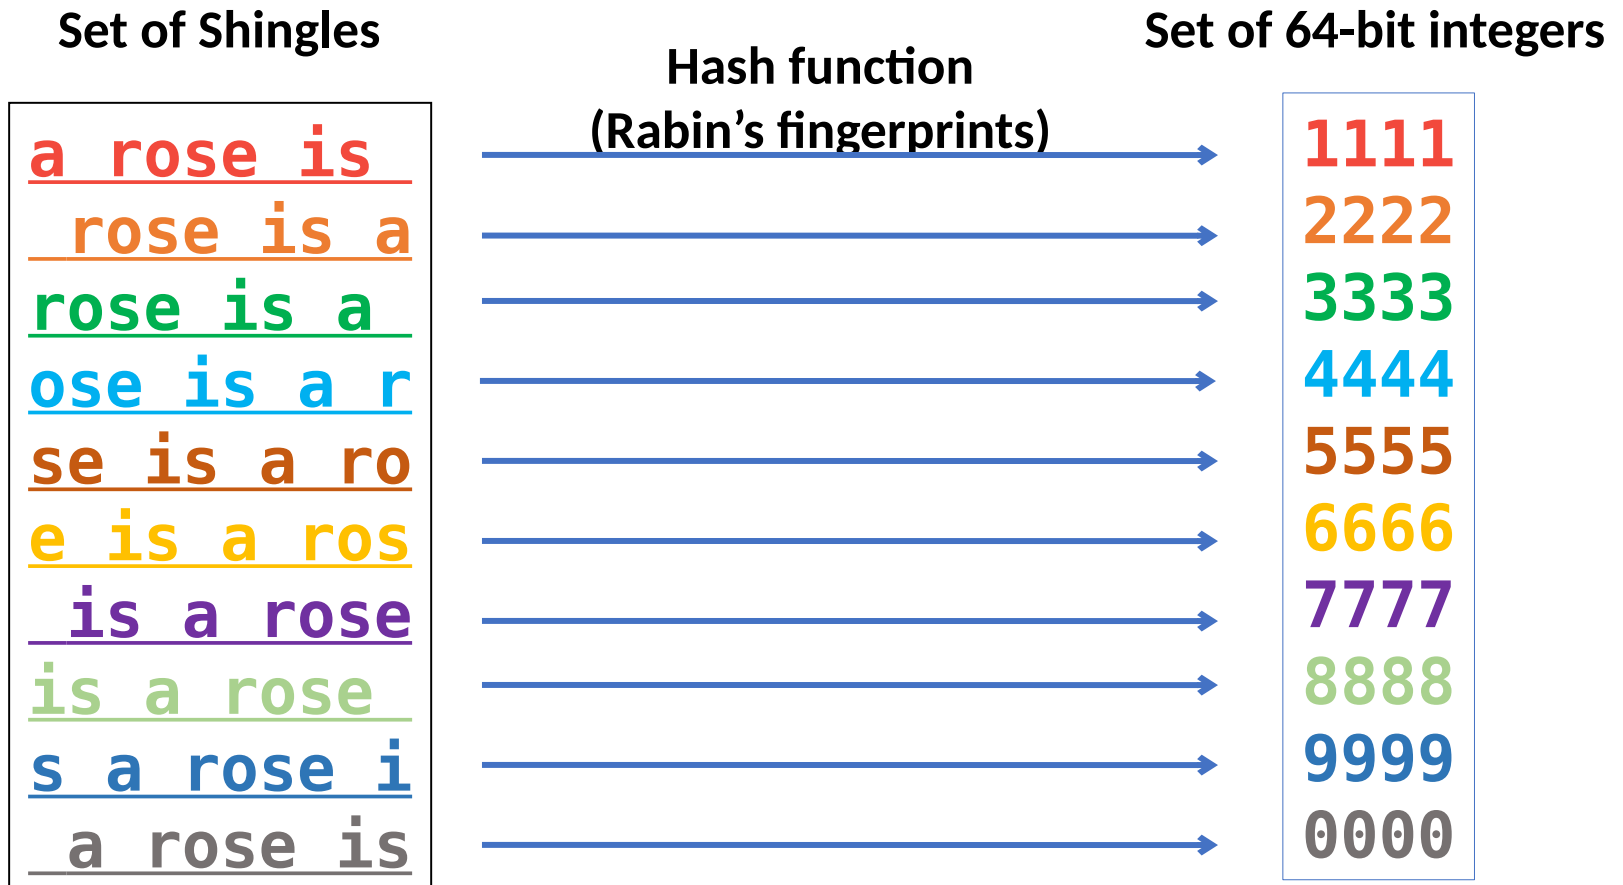
---

- To compress long shingles, we can hash them to (say) 4 bytes.
- Represent a doc by the set of **hash values** of its  $k$ -shingles.
- From now on we will assume that shingles are integers
  - Collisions are possible, but very rare

# Fingerprinting

---

- Hash shingles to 64-bit integers



# Basic Data Model: Sets

---

- **Document**: A document is represented as a **set** shingles (more accurately, hashes of shingles)
- **Document similarity**: **Jaccard** similarity of the sets of shingles.
  - Common shingles over the union of shingles
  - $Sim(C_1, C_2) = |C_1 \cap C_2| / |C_1 \cup C_2|$ .
- Although we use the documents as our driving example the techniques we will describe apply to any kind of sets.
  - E.g., similar customers or items.

# Signatures

---

- **Problem:** shingle sets are too large to be kept in memory.
- **Key idea:** “hash” each set  $S$  to a small **signature**  $\text{Sig}(S)$ , such that:
  1.  $\text{Sig}(S)$  is **small enough** that we can fit a signature in main memory for each set.
  2.  $\text{Sim}(S_1, S_2)$  is (**almost**) the **same** as the “similarity” of  $\text{Sig}(S_1)$  and  $\text{Sig}(S_2)$ . (signature **preserves** similarity).
- **Warning:** This method can produce **false negatives**, and **false positives** (if an additional check is not made).
  - **False negatives:** Similar items deemed as non-similar
  - **False positives:** Non-similar items deemed as similar



# From Sets to Boolean Matrices

---

- Represent the data as a boolean matrix  $M$ 
  - Rows = the universe of all possible set elements
    - In our case, shingle fingerprints take values in  $[0 \dots 2^{64} - 1]$
  - Columns = the sets
    - In our case, documents, sets of shingle fingerprints
  - $M(r, S) = 1$  in row  $r$  and column  $S$ , if and only if  $r$  is a member of  $S$ .
- Typical matrix is sparse.
  - We do not really materialize the matrix

# Example

---

- Universe:  $U = \{A, B, C, D, E, F, G\}$

- $X = \{A, B, F, G\}$

- $Y = \{A, E, F, G\}$

- $\text{Sim}(X, Y) =$

	X	Y
A	1	1
B	1	0
C	0	0
D	0	0
E	0	1
F	1	1
G	1	1

# Example

---

- Universe:  $U = \{A, B, C, D, E, F, G\}$

- $X = \{A, B, F, G\}$

- $Y = \{A, E, F, G\}$

- $\text{Sim}(X, Y) =$

	X	Y
A	1	1
B	1	0
C	0	0
D	0	0
E	0	1
F	1	1
G	1	1

At least one of the columns has value 1

# Example

---

- Universe:  $U = \{A, B, C, D, E, F, G\}$

- $X = \{A, B, F, G\}$

- $Y = \{A, E, F, G\}$

- $\text{Sim}(X, Y) =$

	X	Y
A	1	1
B	1	0
C	0	0
D	0	0
E	0	1
F	1	1
G	1	1

Both columns have value 1

# Minhashing


---

- Pick a **random permutation** of the rows (the universe  $U$ ).
- Define “**hash**” function for set  $S$ 
  - $h(S)$  = the **index** of the **first row** (in the **permuted order**) in which column  $S$  has 1.
  - OR
  - $h(S)$  = the **index** of the **first element** of  $S$  in the **permuted order**.
- Use  $k$  (e.g.,  $k = 100$ ) independent random permutations to create a signature.


# Example of minhash signatures

- Input matrix

	$S_1$	$S_2$	$S_3$	$S_4$
A	1	0	1	0
B	1	0	0	1
C	0	1	0	1
D	0	1	0	1
E	0	1	0	1
F	1	0	1	0
G	1	0	1	0



A
C
G
F
B
E
D




		$S_1$	$S_2$	$S_3$	$S_4$
1	A	1	0	1	0
2	C	0	1	0	1
3	G	1	0	1	0
4	F	1	0	1	0
5	B	1	0	0	1
6	E	0	1	0	1
7	D	0	1	0	1
		1	2	1	2


# Example of minhash signatures

- Input matrix

	$S_1$	$S_2$	$S_3$	$S_4$
A	1	0	1	0
B	1	0	0	1
C	0	1	0	1
D	0	1	0	1
E	0	1	0	1
F	1	0	1	0
G	1	0	1	0



D
B
A
C
F
G
E




		$S_1$	$S_2$	$S_3$	$S_4$
1	D	0	1	0	1
2	B	1	0	0	1
3	A	1	0	1	0
4	C	0	1	0	1
5	F	1	0	1	0
6	G	1	0	1	0
7	E	0	1	0	1
		2	1	3	1


# Example of minhash signatures

- Input matrix

	$S_1$	$S_2$	$S_3$	$S_4$
A	1	0	1	0
B	1	0	0	1
C	0	1	0	1
D	0	1	0	1
E	0	1	0	1
F	1	0	1	0
G	1	0	1	0



C
D
G
F
A
B
E



		$S_1$	$S_2$	$S_3$	$S_4$
1	C	0	1	0	1
2	D	0	1	0	1
3	G	1	0	1	0
4	F	1	0	1	0
5	A	1	0	1	0
6	B	1	0	0	1
7	E	0	1	0	1

3	1	3	1
---	---	---	---



# Example of minhash signatures

- Input matrix

	$S_1$	$S_2$	$S_3$	$S_4$
A	1	0	1	0
B	1	0	0	1
C	0	1	0	1
D	0	1	0	1
E	0	1	0	1
F	1	0	1	0
G	1	0	1	0

$\approx$

Signature matrix

	$S_1$	$S_2$	$S_3$	$S_4$
$h_1$	1	2	1	2
$h_2$	2	1	3	1
$h_3$	3	1	3	1

- $\text{Sig}(S)$  = vector of hash values
  - e.g.,  $\text{Sig}(S_2) = [2, 1, 1]$
- $\text{Sig}(S, i)$  = value of the  $i$ -th hash function for set  $S$ 
  - E.g.,  $\text{Sig}(S_2, 3) = 1$

# Hash function Property

---

$$\Pr(h(S_1) = h(S_2)) = \text{Sim}(S_1, S_2)$$

- where the probability is over all choices of permutations.
- Why?
  - The first row where one of the two sets has value 1 belongs to the union.
    - Recall that union contains rows with at least one 1.
  - We have equality if both sets have value 1, and this row belongs to the intersection


# Example

- Universe:  $U = \{A, B, C, D, E, F, G\}$
- $X = \{A, B, F, G\}$
- $Y = \{A, E, F, G\}$


- Union =  
 $\{A, B, E, F, G\}$
- Intersection =  
 $\{A, F, G\}$

Rows C,D could be anywhere they do not affect the probability

	X	Y
A	1	1
B	1	0
C	0	0
D	0	0
E	0	1
F	1	1
G	1	1



D
*
*
C
*
*
*



	X	Y
D	0	0
C	0	0


# Example

- Universe:  $U = \{A, B, C, D, E, F, G\}$
- $X = \{A, B, F, G\}$
- $Y = \{A, E, F, G\}$


The \* rows belong to the union

- Union =  
 $\{A, B, E, F, G\}$
- Intersection =  
 $\{A, F, G\}$

	X	Y
A	1	1
B	1	0
C	0	0
D	0	0
E	0	1
F	1	1
G	1	1



D
*
*
C
*
*
*



	X	Y
D	0	0
C	0	0


# Example

- Universe:  $U = \{A, B, C, D, E, F, G\}$
- $X = \{A, B, F, G\}$
- $Y = \{A, E, F, G\}$


- Union =  
 $\{A, B, E, F, G\}$
- Intersection =  
 $\{A, F, G\}$

The question is what is the value of the **first** \* element

	X	Y
A	1	1
B	1	0
C	0	0
D	0	0
E	0	1
F	1	1
G	1	1



D
*
*
C
*
*
*



	X	Y
D	0	0
C	0	0


# Example

- Universe:  $U = \{A, B, C, D, E, F, G\}$
- $X = \{A, B, F, G\}$
- $Y = \{A, E, F, G\}$


If it belongs to the intersection, then  
 $h(X) = h(Y)$

- Union =  
 $\{A, B, E, F, G\}$
- Intersection =  
 $\{A, F, G\}$

	X	Y
A	1	1
B	1	0
C	0	0
D	0	0
E	0	1
F	1	1
G	1	1



D
*
*
C
*
*
*



	X	Y
D	0	0
C	0	0

# Example


- Universe:  $U = \{A, B, C, D, E, F, G\}$
- $X = \{A, B, F, G\}$
- $Y = \{A, E, F, G\}$

Every element of the union is equally likely to be the \* element


$$\Pr(h(X) = h(Y)) = \text{Sim}(X, Y)$$

- Union =  
 $\{A, B, E, F, G\}$
- Intersection =  
 $\{A, F, G\}$

	X	Y
A	1	1
B	1	0
C	0	0
D	0	0
E	0	1
F	1	1
G	1	1



D
*
*
C
*
*
*



	X	Y
D	0	0
C	0	0

# Similarity for Signatures

- The **similarity of signatures** is the **fraction of the hash functions** in which they agree.

	$S_1$	$S_2$	$S_3$	$S_4$
A	1	0	1	0
B	1	0	0	1
C	0	1	0	1
D	0	1	0	1
E	0	1	0	1
F	1	0	1	0
G	1	0	1	0

$\approx$

Signature matrix

$S_1$	$S_2$	$S_3$	$S_4$
1	2	1	2
2	1	3	1
3	1	3	1

	Actual	Sig
$(S_1, S_2)$	0	0
$(S_1, S_3)$	3/5	2/3
$(S_1, S_4)$	1/7	0
$(S_2, S_3)$	0	0
$(S_2, S_4)$	3/4	1
$(S_3, S_4)$	0	0

Zero similarity is preserved  
High similarity is well approximated



# Is it now feasible?

---

- Assume a billion rows
- Hard to pick a random permutation of 1 billion...
- **Even representing a random permutation requires 1 billion entries!!!**
- How about accessing rows in permuted order?



# Being more practical

---

- Instead of permuting the rows we will apply a **hash function** that maps the rows to a new (possibly larger) space
  - The value of the hash function is the position of the row in the new order (permutation).
  - Each set is represented by the smallest hash value among the elements in the set
- The space of the hash functions should be such that if we select one at random each element (row) has equal probability to have the smallest value
  - **Min-wise independent** hash functions

# Algorithm – One set, one hash function

Computing  $\text{Sig}(\mathbf{S}, i)$  for a single column  $\mathbf{S}$  and single hash function  $h_i$

for each row  $r$

In practice only the rows (shingles) that appear in the data

compute  $h_i(r)$

$h_i(r)$  = index of row  $r$  in permutation

if column  $\mathbf{S}$  that has 1 in row  $r$

$\mathbf{S}$  contains row  $r$

if  $h_i(r)$  is a smaller value than  $\text{Sig}(\mathbf{S}, i)$

then

Find the row  $r$  with minimum index

$\text{Sig}(\mathbf{S}, i) = h_i(r)$

$\text{Sig}(\mathbf{S}, i)$  will become the smallest value of  $h_i(r)$  among all rows (shingles) for which column  $\mathbf{S}$  has value 1 (shingle belongs in  $\mathbf{S}$ ); i.e.,  $h_i(r)$  gives the min index for the  $i$ -th permutation

# Algorithm – All sets, $k$ hash functions

---

Pick  $k=100$  hash functions  $(h_1, \dots, h_k)$

In practice this means selecting the hash function parameters

**for** each row  $r$

**for** each hash function  $h_i$

compute  $h_i(r)$

Compute  $h_i(r)$  only once for all sets

**for** each column  $S$  that has  $1$  in row  $r$

**if**  $h_i(r)$  is a smaller value than  $\text{Sig}(S,i)$   
**then**

$\text{Sig}(S,i) = h_i(r);$

# Algorithm – All sets, **k** hash functions

x	Row	S1	S2	h(x)	g(x)		Sig1	Sig2
0	A	1	0	1	3	$h(0) = 1$	1	-
1	B	0	1	2	0	$g(0) = 3$	3	-
2	C	1	1	3	2	$h(1) = 2$	1	2
3	D	1	0	4	4	$g(1) = 0$	3	0
4	E	0	1	0	1			

$$h(x) = x+1 \bmod 5$$

$$g(x) = 2x+3 \bmod 5$$

h(Row)	Row	S1	S2	g(Row)	Row	S1	S2
0	E	0	1	0	B	0	1
1	A	1	0	1	E	0	1
2	B	0	1	2	C	1	0
3	C	1	1	3	A	1	1
4	D	1	0	4	D	1	0

$h(2) = 3$	1	2
$g(2) = 2$	2	0
$h(3) = 4$	1	2
$g(3) = 4$	2	0
$h(4) = 0$	1	0
$g(4) = 1$	2	0

# Implementation

---

- Often, data is given by column, not row.
  - E.g., columns = documents, rows = shingles.
- If so, sort matrix once so it is by row.
- And **always** compute  $h_i(r)$  only once for each row.

# Finding similar pairs

---

- Problem: Find all pairs of documents with similarity at least  $t = 0.8$
- While the signatures of all columns may fit in main memory, comparing the signatures of all pairs of columns is **quadratic** in the number of columns.
- **Example:**  $10^6$  columns implies  $5 \cdot 10^{11}$  column-comparisons.
- At 1 microsecond/comparison: 6 days.

# Locality-Sensitive Hashing

---


- **What we want:** a function  $f(X,Y)$  that tells whether or not  $X$  and  $Y$  is a **candidate pair**: a pair of elements whose similarity must be evaluated.
- **A simple idea:**  $X$  and  $Y$  are a candidate pair if they have the **same min-hash signature**.
  - Easy to test by **hashing** the **signatures**.
  - **Similar sets** are more **likely** to have the **same signature**.
  - Likely to produce many **false negatives**.

! Multiple levels of Hashing!

    - Requiring full match of signature is strict, some similar sets will be lost.
- **Improvement:** Compute multiple signatures; candidate pairs should have **at least** one common signature.
  - Reduce the probability for false negatives.



# Signature matrix reminder



# Partition into Bands – (1)


---

- Divide the signature matrix Sig into  $b$  bands of  $r$  rows.
  - Each band is a mini-signature with  $r$  hash functions.

# Partition into Bands – (1)

---


$n = b * r$  hash functions



# Partition into Bands – (2)

---

- Divide the signature matrix Sig into  $b$  bands of  $r$  rows.
  - Each band is a mini-signature with  $r$  hash functions.
- For each band, hash the mini-signature to a hash table with  $k$  buckets.
  - Make  $k$  as large as possible so that mini-signatures that hash to the same bucket are almost certainly identical.




# Partition into Bands – (3)

---

- Divide the signature matrix Sig into  $b$  bands of  $r$  rows.
  - Each band is a **mini-signature** with  $r$  hash functions.
- For each band, hash the mini-signature to a hash table with  $k$  buckets.
  - Make  $k$  as large as possible so that mini-signatures that hash to the same bucket are **almost certainly identical**.
- **Candidate** column pairs are those that hash to the same bucket for **at least** 1 band.
- Tune  $b$  and  $r$  to catch **most similar pairs**, but **few non-similar pairs**.


# Analysis of LSH – What we want

---




# What One Band of One Row Gives You

---





# What $b$ Bands of $r$ Rows Gives You



# Example: $b = 20$ ; $r = 5$

---

$s$	$1-(1-s^r)^b$
.2	.006
.3	.047
.4	.186
.5	.470
.6	.802
.7	.975
.8	.9996




Figure 3.7: The S-curve

# Suppose $S_1, S_2$ are 80% Similar

---

- We want all 80%-similar pairs. Choose 20 bands of 5 integers/band.
- Probability  $S_1, S_2$  identical in one particular band:  
$$(0.8)^5 = 0.328.$$
- Probability  $S_1, S_2$  are not similar in any of the 20 bands:  
$$(1-0.328)^{20} = 0.00035$$

i.e., about 1/3000-th of the 80%-similar column pairs are false negatives.

- Probability  $S_1, S_2$  are similar in at least one of the 20 bands:  
$$1-0.00035 = 0.999$$

# Suppose $S_1, S_2$ Only 40% Similar

---

- Probability  $S_1, S_2$  identical in any one particular band:  
 $(0.4)^5 = 0.01$  .
- Probability  $S_1, S_2$  identical in **at least** 1 of 20 bands:  
 $\leq 20 * 0.01 = 0.2$  .
- But **false positives** much lower for similarities  $\ll 40\%$ .