

CHAPTER - 6

Basic SQL

Note: Slides, content, web links and end chapter questions are prepared from Pearson textbook (Elmasri & Navathe), and other Internet resources.

Topics of Discussion

- A. Introduction to Structure Query Language.
- B. SQL Data Definition and Data Types
 - A. CREATE SCHEMA
 - B. CREATE TABLE
- C. Specifying Constraints in SQL
- D. INSERT, DELETE, and UPDATE Statements in SQL
- E. Basic Retrieval Queries in SQL
- F. Additional Features of SQL

Introduction - Basic SQL

- Structure Query Language (SQL)

- Considered one of the major reasons for the commercial success of relational databases.

- SQL

- The origin of SQL is relational predicate calculus called tuple calculus, which was proposed initially as the language SQUARE.
 - (SQL Actually comes from the word “SEQUEL” which was the original term used in the research article: “SEQUEL TO SQUARE” by Chamberlin and Boyce. IBM could not copyright that term, so they abbreviated it to SQL and copyrighted the term SQL.)
 - Now popularly known as “Structured Query language”.
 - SQL is an informal or practical rendering of the relational data model with syntax.

Introduction - Basic SQL

- SQL has gone through many standards: starting with SQL-86 or SQL1. SQL-92 is referred to as SQL-2.
- Later standards (from SQL-1999) are divided into **core** specifications and specialized **extensions**. The extensions are implemented for different applications – such as data mining, data warehousing, multimedia, etc.
- SQL-2006 added XML features; In 2008, added Object-Oriented features.
- SQL-3 is the current standard which started with SQL-1999 and latest within SQL-3 is SQL-2022.
- It is not fully implemented in any (mostly) RDBMS.

Introduction - Basic SQL

- We will discuss the basic standard SQL syntax – there are variations in existing RDBMS systems.
- For example, there may little variation of SQL used for Oracle and for MS SQL Server.
- **SQL schema**
 - Identified by a **schema name**
 - Includes an **authorization identifier** and **descriptors** for each element
 - **Schema elements** include
 - Tables, constraints, views, domains, and other constructs

SQL Data Definition, Data Types

- Terminology:
 - **Table**, **row**, and **column** used for relational model terms **relation**, **tuple**, and **attribute**.
- CREATE statement
 - Main SQL command for data definition
- CREATE SCHEMA statement
 - CREATE SCHEMA **COMPANY** AUTHORIZATION 'apatel38';
- Catalog
 - Named collection of schemas in an SQL environment
- SQL also has the concept of a cluster of catalogs.
- Each statement in SQL ends with a **semicolon**

The CREATE TABLE Command in SQL

- Specifying a new relation (table)
 - Provide name of table
 - Specify attributes, their types and initial constraints
- Can optionally specify schema:
 - `CREATE TABLE COMPANY.EMPLOYEE`
or
 - `CREATE TABLE EMPLOYEE`
- Base tables (base relations)
 - Relation and its tuples are actually created and stored as a file by the DBMS. (Internally when it stored at physical storage like hard disk, it is known as **flat file**.)
- Virtual relations (views)
 - Created through the `CREATE VIEW` statement.
 - Do not correspond to any physical file.

Example : database schema (Fig. 5.7)

EMPLOYEE

| Fname | Minit | Lname | <u>Ssn</u> | Bdate | Address | Sex | Salary | Super_ssn | Dno |
|-------|-------|-------|------------|-------|---------|-----|--------|-----------|-----|
|-------|-------|-------|------------|-------|---------|-----|--------|-----------|-----|

DEPARTMENT

| Dname | <u>Dnumber</u> | Mgr_ssn | Mgr_start_date |
|-------|----------------|---------|----------------|
|-------|----------------|---------|----------------|

DEPT_LOCATIONS

| <u>Dnumber</u> | <u>Dlocation</u> |
|----------------|------------------|
|----------------|------------------|

PROJECT

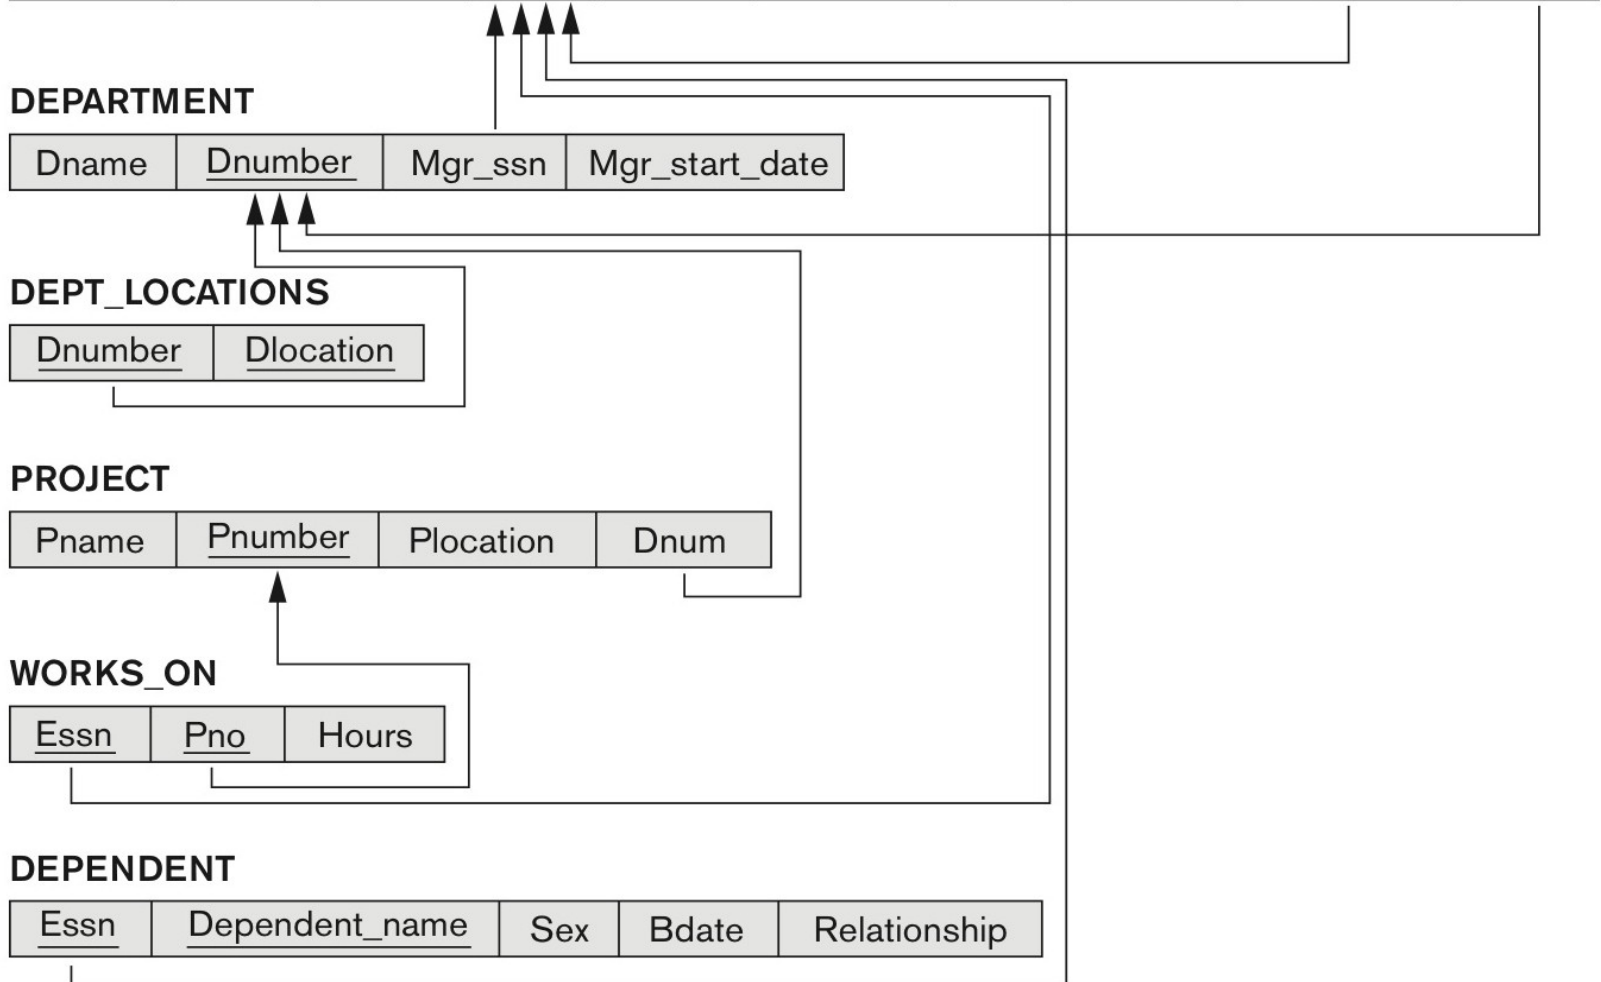
| Pname | <u>Pnumber</u> | Plocation | Dnum |
|-------|----------------|-----------|------|
|-------|----------------|-----------|------|

WORKS_ON

| <u>Essn</u> | <u>Pno</u> | Hours |
|-------------|------------|-------|
|-------------|------------|-------|

DEPENDENT

| <u>Essn</u> | <u>Dependent_name</u> | Sex | Bdate | Relationship |
|-------------|-----------------------|-----|-------|--------------|
|-------------|-----------------------|-----|-------|--------------|



Example : database state (Fig. 5.6)

EMPLOYEE

| Fname | Minit | Lname | <u>Ssn</u> | Bdate | Address | Sex | Salary | Super_ssn | Dno |
|----------|-------|---------|------------|------------|--------------------------|-----|--------|-----------|-----|
| John | B | Smith | 123456789 | 1965-01-09 | 731 Fondren, Houston, TX | M | 30000 | 333445555 | 5 |
| Franklin | T | Wong | 333445555 | 1955-12-08 | 638 Voss, Houston, TX | M | 40000 | 888665555 | 5 |
| Alicia | J | Zelaya | 999887777 | 1968-01-19 | 3321 Castle, Spring, TX | F | 25000 | 987654321 | 4 |
| Jennifer | S | Wallace | 987654321 | 1941-06-20 | 291 Berry, Bellaire, TX | F | 43000 | 888665555 | 4 |
| Ramesh | K | Narayan | 666884444 | 1962-09-15 | 975 Fire Oak, Humble, TX | M | 38000 | 333445555 | 5 |
| Joyce | A | English | 453453453 | 1972-07-31 | 5631 Rice, Houston, TX | F | 25000 | 333445555 | 5 |
| Ahmad | V | Jabbar | 987987987 | 1969-03-29 | 980 Dallas, Houston, TX | M | 25000 | 987654321 | 4 |
| James | E | Borg | 888665555 | 1937-11-10 | 450 Stone, Houston, TX | M | 55000 | NULL | 1 |

DEPARTMENT

| Dname | <u>Dnumber</u> | Mgr_ssn | Mgr_start_date |
|----------------|----------------|-----------|----------------|
| Research | 5 | 333445555 | 1988-05-22 |
| Administration | 4 | 987654321 | 1995-01-01 |
| Headquarters | 1 | 888665555 | 1981-06-19 |

DEPT_LOCATIONS

| <u>Dnumber</u> | <u>Dlocation</u> |
|----------------|------------------|
| 1 | Houston |
| 4 | Stafford |
| 5 | Bellaire |
| 5 | Sugarland |
| 5 | Houston |

continued on next slide

Example : database state (Fig. 5.6)

WORKS_ON

| <u>Essn</u> | <u>Pno</u> | Hours |
|-------------|------------|-------|
| 123456789 | 1 | 32.5 |
| 123456789 | 2 | 7.5 |
| 666884444 | 3 | 40.0 |
| 453453453 | 1 | 20.0 |
| 453453453 | 2 | 20.0 |
| 333445555 | 2 | 10.0 |
| 333445555 | 3 | 10.0 |
| 333445555 | 10 | 10.0 |
| 333445555 | 20 | 10.0 |
| 999887777 | 30 | 30.0 |
| 999887777 | 10 | 10.0 |
| 987987987 | 10 | 35.0 |
| 987987987 | 30 | 5.0 |
| 987654321 | 30 | 20.0 |
| 987654321 | 20 | 15.0 |
| 888665555 | 20 | NULL |

PROJECT

| Pname | <u>Pnumber</u> | Plocation | Dnum |
|-----------------|----------------|-----------|------|
| ProductX | 1 | Bellaire | 5 |
| ProductY | 2 | Sugarland | 5 |
| ProductZ | 3 | Houston | 5 |
| Computerization | 10 | Stafford | 4 |
| Reorganization | 20 | Houston | 1 |
| Newbenefits | 30 | Stafford | 4 |

DEPENDENT

| <u>Essn</u> | <u>Dependent_name</u> | Sex | Bdate | Relationship |
|-------------|-----------------------|-----|------------|--------------|
| 333445555 | Alice | F | 1986-04-05 | Daughter |
| 333445555 | Theodore | M | 1983-10-25 | Son |
| 333445555 | Joy | F | 1958-05-03 | Spouse |
| 987654321 | Abner | M | 1942-02-28 | Spouse |
| 123456789 | Michael | M | 1988-01-04 | Son |
| 123456789 | Alice | F | 1988-12-30 | Daughter |
| 123456789 | Elizabeth | F | 1967-05-05 | Spouse |

CREATE TABLE data definition statements for defining the COMPANY schema (Figure 6.1)

```
CREATE TABLE EMPLOYEE
  ( Fname          VARCHAR(15)          NOT NULL,
    Minit          CHAR,
    Lname          VARCHAR(15)          NOT NULL,
    Ssn            CHAR(9)              NOT NULL,
    Bdate          DATE,
    Address        VARCHAR(30),
    Sex            CHAR,
    Salary         DECIMAL(10,2),
    Super_ssn      CHAR(9),
    Dno            INT                  NOT NULL,
    PRIMARY KEY (Ssn),
CREATE TABLE DEPARTMENT
  ( Dname          VARCHAR(15)          NOT NULL,
    Dnumber        INT                  NOT NULL,
    Mgr_ssn        CHAR(9)              NOT NULL,
    Mgr_start_date DATE,
    PRIMARY KEY (Dnumber),
    UNIQUE (Dname),
    FOREIGN KEY (Mgr_ssn) REFERENCES EMPLOYEE(Ssn) );
CREATE TABLE DEPT_LOCATIONS
  ( Dnumber        INT                  NOT NULL,
    Dlocation      VARCHAR(15)          NOT NULL,
    PRIMARY KEY (Dnumber, Dlocation),
    FOREIGN KEY (Dnumber) REFERENCES DEPARTMENT(Dnumber) );
```

continued next slide

CREATE TABLE data definition statements for defining the COMPANY schema (Figure 6.1)

CREATE TABLE PROJECT

| | | |
|-----------|--------------|-----------|
| (Pname | VARCHAR(15) | NOT NULL, |
| Pnumber | INT | NOT NULL, |
| Plocation | VARCHAR(15), | |
| Dnum | INT | NOT NULL, |

PRIMARY KEY (Pnumber),
UNIQUE (Pname),
FOREIGN KEY (Dnum) **REFERENCES** DEPARTMENT(Dnumber));

CREATE TABLE WORKS_ON

| | | |
|--------|--------------|-----------|
| (Essn | CHAR(9) | NOT NULL, |
| Pno | INT | NOT NULL, |
| Hours | DECIMAL(3,1) | NOT NULL, |

PRIMARY KEY (Essn, Pno),
FOREIGN KEY (Essn) **REFERENCES** EMPLOYEE(Ssn),
FOREIGN KEY (Pno) **REFERENCES** PROJECT(Pnumber));

CREATE TABLE DEPENDENT

| | | |
|----------------|-------------|-----------|
| (Essn | CHAR(9) | NOT NULL, |
| Dependent_name | VARCHAR(15) | NOT NULL, |
| Sex | CHAR, | |
| Bdate | DATE, | |
| Relationship | VARCHAR(8), | |

PRIMARY KEY (Essn, Dependent_name),
FOREIGN KEY (Essn) **REFERENCES** EMPLOYEE(Ssn));

Attribute Data Types and Domains in SQL

- **Basic data types**

- **Numeric data types**

- Integer numbers: INTEGER, INT, and SMALLINT
 - Floating-point (real) numbers: FLOAT or REAL, and DOUBLE PRECISION

- **Character-string data types**

- Fixed length: CHAR(*n*), CHARACTER(*n*)
 - Varying length: VARCHAR(*n*), CHAR VARYING(*n*), CHARACTER VARYING(*n*)

- **Boolean data type**

- Values of TRUE or FALSE or NULL

Attribute Data Types and Domains in SQL

- **DATE** data type
 - Ten positions. Components are YEAR, MONTH, and DAY in the form YYYY-MM-DD
 - Multiple mapping functions available in RDBMSs to change date formats
- **Timestamp** data type Includes the DATE and TIME fields
 - Plus a minimum of six positions for decimal fractions of seconds
 - Optional WITH TIME ZONE qualifier
- **INTERVAL** data type
 - Specifies a relative value that can be used to increment or decrement an absolute value of a date, time, or timestamp

Attribute Data Types and Domains in SQL

- **DATE, TIME, Timestamp, INTERVAL** data types can be **cast** or converted to string formats for comparison **DATE** data type.
- **Bit-string** data types
 - Fixed length: `BIT (n)` Varying length: `BIT VARYING (n)` .
 - Literal bit strings example `B'10101'` .
- **BLOB (Binary Large Object)**
 - For large binary value such as images.

Attribute Data Types and Domains in SQL

■ Domain

- Name used with the attribute specification
- Makes it easier to change the data type for a domain that is used by numerous attributes
- Improves schema readability
- Example:
 - `CREATE DOMAIN SSN_TYPE AS CHAR(9);`

■ TYPE

- User Defined Types (UDTs) are supported for object-oriented applications.
- Uses the command: `CREATE TYPE`

Specifying Constraints in SQL

Basic constraints:

- Relational Model has 3 basic constraint types that are supported in SQL:
 - **Key** constraint:
A primary key value cannot be duplicated
 - **Entity Integrity** Constraint:
A “primary key” value cannot be null
 - **Referential integrity** constraints :
The “foreign key “ must have a value that is already present as a primary key or may be null.

Specifying Attribute Constraints

- Other Restrictions on attribute domains:
- Default value of an attribute
 - **DEFAULT** <value>
- **NULL** is not permitted for a particular attribute (**NOT NULL**)
- **CHECK** clause
 - `Dnumber INT NOT NULL CHECK (Dnumber > 0 AND Dnumber < 10);`

Specifying Key and Referential Integrity Constraints

- **PRIMARY KEY** clause

- Specifies one or more attributes that make up the primary key of a relation
- `Dnumber INT PRIMARY KEY;`

- **UNIQUE** clause

- Specifies alternate (secondary) keys (called **CANDIDATE** keys in the relational model).
- `Dname VARCHAR(15) UNIQUE;`

Specifying Key and Referential Integrity Constraints

- **FOREIGN KEY** clause
 - Default operation: reject update on violation
 - Attach **referential triggered action** clause
 - Options include SET NULL, CASCADE, and SET DEFAULT
 - Action taken by the DBMS for SET NULL or SET DEFAULT is the same for both ON DELETE and ON UPDATE
 - CASCADE option suitable for “relationship” relations
- Giving Name to Constraint
 - Using the Keyword **CONSTRAINT**
 - Name a constraint. Useful for later altering

Default attribute values and referential integrity triggered action specification in SQL (Figure 6.2)

```
CREATE TABLE EMPLOYEE
( ... ,
  Dno          INT          NOT NULL          DEFAULT 1,
  CONSTRAINT EMPPK
    PRIMARY KEY (Ssn),
  CONSTRAINT EMPSUPERFK
    FOREIGN KEY (Super_ssn) REFERENCES EMPLOYEE(Ssn)
      ON DELETE SET NULL          ON UPDATE CASCADE,
  CONSTRAINT EMPDEPTFK
    FOREIGN KEY(Dno) REFERENCES DEPARTMENT(Dnumber)
      ON DELETE SET DEFAULT      ON UPDATE CASCADE);

CREATE TABLE DEPARTMENT
( ... ,
  Mgr_ssn CHAR(9)          NOT NULL          DEFAULT '888665555',
  ... ,
  CONSTRAINT DEPTPK
    PRIMARY KEY(Dnumber),
  CONSTRAINT DEPTSK
    UNIQUE (Dname),
  CONSTRAINT DEPTMGRFK
    FOREIGN KEY (Mgr_ssn) REFERENCES EMPLOYEE(Ssn)
      ON DELETE SET DEFAULT      ON UPDATE CASCADE);

CREATE TABLE DEPT_LOCATIONS
( ... ,
  PRIMARY KEY (Dnumber, Dlocation),
  FOREIGN KEY (Dnumber) REFERENCES DEPARTMENT(Dnumber)
      ON DELETE CASCADE          ON UPDATE CASCADE);
```

Specifying Constraints on Tuples Using CHECK

```
CONSTRAINT EMPSUPERFK  
FOREIGN KEY (Super_ssn) REFERENCES EMPLOYEE(Ssn)  
ON DELETE SET NULL ON UPDATE CASCADE,
```

- This means that if the row for supervising employee is deleted, the value of Super_ssn is automatically set to NULL for all employee rows that were referencing the deleted employee row.
- On the other hand, if the Ssn value for a supervising employee is updated, (in case of incorrect entry of Ssn) the new value is cascaded to super_ssn for all employee rows referencing the updated employee row.
- Additional Constraints on individual tuples within a relation are also possible using CHECK
- CHECK clauses at the end of a CREATE TABLE statement
 - Apply to each tuple individually
 - CHECK (Dept_create_date <= Mgr_start_date);

INSERT, DELETE, and UPDATE Statements in SQL

- Three commands used to modify the database:
 - INSERT, DELETE, and UPDATE
- INSERT typically inserts a tuple (row) in a relation (table)
- UPDATE may update a number of tuples (rows) in a relation (table) that satisfy the condition
- DELETE may also update a number of tuples (rows) in a relation (table) that satisfy the condition

INSERT

- In its simplest form, it is used to add one or more rows to a table.
- Attribute values should be listed in the same order as the attributes specified in the CREATE TABLE command
- Constraints on data types are observed automatically
- Any integrity constraints as a part of the DDL specification are enforced

The INSERT Command

- Specify the relation name and a list of values for the tuple. All values including nulls are supplied.

```
U1:  INSERT INTO  EMPLOYEE
      VALUES      ( 'Richard', 'K', 'Marini', '653298653', '1962-12-30', '98
                    Oak Forest, Katy, TX', 'M', 37000, '653298653', 4 );
```

- The variation below inserts multiple tuples where a new table is loaded values from the result of a query.

```
U3B:  INSERT INTO  WORKS_ON_INFO ( Emp_name, Proj_name,
                                     Hours_per_week )
      SELECT        E.Lname, P.Pname, W.Hours
      FROM          PROJECT P, WORKS_ON W, EMPLOYEE E
      WHERE         P.Pnumber=W.Pno AND W.Essn=E.Ssn;
```

BULK LOADING OF TABLES

- Most DBMSs have bulk loading tools that allow a user to load formatted data from a file into a table without having to write a large number of INSERT commands.
- One variation of **INSERT** is used for bulk-loading of several tuples into tables is as follow.
 - Example: A new table D5EMPS can be created with the same attributes as EMPLOYEE and using LIKE and DATA in the syntax, it can be loaded with entire data. (or load it with rows of employees who work in department 5.
 - Can be written as below.

```
CREATE TABLE D5EMPS LIKE EMPLOYEE  
      (SELECT E.*  
      FROM EMPLOYEE AS E  
      WHERE E.Dno=5)  
  
WITH DATA;
```

UPDATE

- Used to modify attribute values of one or more selected tuples
- A WHERE-clause selects the tuples to be modified
- An additional SET-clause specifies the attributes to be modified and their new values
- Each command modifies tuples *in the same relation*
- Referential integrity specified as part of DDL specification is enforced
- Ex1: Change the location and controlling department number of project number 10 to 'Bellaire' and 5, respectively.

```
U5:      UPDATE      PROJECT
          SET          PLOCATION = 'Bellaire', DNUM
          = 5
          WHERE        PNUMBER=10
```

UPDATE

- Ex2: Employee whose ssn is '333445555' is now transferred to department 5.
 - U1: UPDATE Employee
SET dno = 5 WHERE ssn = '333445555';
- Ex3: Give all employees in the 'Research' department a 10% raise in salary.

U6: UPDATE EMPLOYEE
SET SALARY = SALARY * 1.1
WHERE DNO IN (SELECT DNUMBER
FROM DEPARTMENT
WHERE DNAME='Research')

- In this request, the modified SALARY value depends on the original SALARY value in each tuple
 - The reference to the SALARY attribute on the right of = refers to the old SALARY value before modification
 - The reference to the SALARY attribute on the left of = refers to the new SALARY value after modification

DELETE

- Removes tuples from a relation
 - Includes a WHERE-clause to select the tuples to be deleted.
 - Referential integrity should be enforced
 - Tuples are deleted from only *one table* at a time (unless CASCADE is specified on a referential integrity constraint).
 - A missing WHERE-clause specifies that *all tuples* in the relation are to be deleted; the table then becomes an empty table.
 - The number of tuples deleted depends on the number of tuples in the relation that satisfy the WHERE-clause.

The DELETE Command

- Removes rows from a table
 - Includes a `WHERE` clause to select the tuples to be deleted. The number of tuples deleted will vary.

U4A: **DELETE FROM** **EMPLOYEE**
 WHERE **Lname='Brown';**

U4B: **DELETE FROM** **EMPLOYEE**
 WHERE **Ssn='123456789';**

U4C: **DELETE FROM** **EMPLOYEE**
 WHERE **Dno=5;**

U4D: **DELETE FROM** **EMPLOYEE;**

Basic Retrieval Queries in SQL

- **SELECT statement**
 - One basic statement for retrieving information from a database
- SQL allows a table to have two or more tuples that are identical in all their attribute values
 - Unlike relational model (relational model is strictly set-theory based). Multiset or bag behavior
 - Tuple-id may be used as a key

```
SELECT    <attribute list>
FROM      <table list>
WHERE     <condition>;
```

where

- <attribute list> is a list of attribute names whose values are to be retrieved by the query.
- <table list> is a list of the relation names required to process the query.
- <condition> is a conditional (Boolean) expression that identifies the tuples to be retrieved by the query.

Basic Retrieval Queries in SQL

- Logical comparison operators
 - =, <, <=, >, >=, and <>
- **Projection attributes**
 - Attributes whose values are to be retrieved
- **Selection condition**
 - Boolean condition that must be true for any retrieved tuple. Selection conditions include join conditions, when multiple relations are involved.

Basic Retrieval Queries in SQL

| <u>Bdate</u> | <u>Address</u> |
|--------------|--------------------------|
| 1965-01-09 | 731 Fondren, Houston, TX |

| <u>Fname</u> | <u>Lname</u> | <u>Address</u> |
|--------------|--------------|--------------------------|
| John | Smith | 731 Fondren, Houston, TX |
| Franklin | Wong | 638 Voss, Houston, TX |
| Ramesh | Narayan | 975 Fire Oak, Humble, TX |
| Joyce | English | 5631 Rice, Houston, TX |

Query 0. Retrieve the birth date and address of the employee(s) whose name is 'John B. Smith'.

Q0: **SELECT** Bdate, Address
 FROM EMPLOYEE
 WHERE Fname='John' **AND** Minit='B' **AND** Lname='Smith';

Query 1. Retrieve the name and address of all employees who work for the 'Research' department.

Q1: **SELECT** Fname, Lname, Address
 FROM EMPLOYEE, DEPARTMENT
 WHERE Dname='Research' **AND** Dnumber=Dno;

Basic Retrieval Queries in SQL

(c)

| <u>Pnumber</u> | <u>Dnum</u> | <u>Lname</u> | <u>Address</u> | <u>Bdate</u> |
|----------------|-------------|--------------|------------------------|--------------|
| 10 | 4 | Wallace | 291Berry, Bellaire, TX | 1941-06-20 |
| 30 | 4 | Wallace | 291Berry, Bellaire, TX | 1941-06-20 |

Query 2. For every project located in 'Stafford', list the project number, the controlling department number, and the department manager's last name, address, and birth date.

Q2: **SELECT** Pnumber, Dnum, Lname, Address, Bdate
 FROM PROJECT, DEPARTMENT, EMPLOYEE
 WHERE Dnum=Dnumber **AND** Mgr_ssn=Ssn **AND**
 Plocation='Stafford';

Ambiguous Attribute Names

- Same name can be used for two (or more) attributes in different relations (tables)
 - As long as the attributes are in different tables.
 - Must **qualify** the attribute name with the table name to prevent ambiguity.
 - Ex: Dno attribute of the EMPLOYEE table was called Dnumber and Dname attribute of DEPARTMENT was also called Name; to prevent ambiguity, query Q1 would be rephrased as shown in Q1A, as below.

Q1A: `select fname, Lname, address
from dbo.employee, dbo.department
where Department.Dname = 'Research' and Department.Dnumber = Employee.Dno;`

Aliasing, and Renaming

- **Aliases or tuple variables**

- Declare alternative relation names E and S to refer to the EMPLOYEE relation twice in a query:

Q8. For each employee, retrieve the employee's first and last name and the first and last name of his or her immediate supervisor.

- **SELECT** E.Fname, E.Lname, S.Fname, S.Lname
FROM EMPLOYEE **AS** E, EMPLOYEE **AS** S
WHERE E.Super_ssn=S.Ssn;
- Recommended practice to abbreviate names and to prefix same or similar attribute from multiple tables.

Aliasing, Renaming and Tuple Variables

- The attribute names can also be renamed

```
EMPLOYEE AS E (Fn, Mi, Ln, Ssn, Bd, Addr,  
Sex, Sal, Sssn, Dno)
```

- Note that the table EMPLOYEE now has a variable name E which corresponds to a tuple variable
- The “AS” may be dropped in most SQL implementations

Unspecified WHERE Clause and Use of *

- Missing WHERE clause
 - Indicates no condition on tuple selection
- Effect is a CROSS PRODUCT
 - Result is all possible tuple combinations (or the Algebra operation of Cartesian Product) result

Queries 9 and 10. Select all EMPLOYEE Ssns (Q9) and all combinations of EMPLOYEE Ssn and DEPARTMENT Dname (Q10) in the database.

Q9: **SELECT** Ssn
 FROM EMPLOYEE;

Q10: **SELECT** Ssn, Dname
 FROM EMPLOYEE, DEPARTMENT;

Unspecified WHERE Clause and Use of *

- Specify an asterisk (*)
 - Retrieve all the attribute values of the selected tuples
 - The * can be prefixed by the relation name; e.g., EMPLOYEE *
 - Query Q1C retrieves all the attribute values of any EMPLOYEE who works in DEPARTMENT number 5 (Figure 6.3(g))

Q1C: **SELECT** *

FROM EMPLOYEE

WHERE Dno=5;

Figure 6.3g Results of SQL queries when applied to the COMPANY database state shown in Figure 5.6. Q1C.

| <u>Fname</u> | <u>Minit</u> | <u>Lname</u> | Ssn | <u>Bdate</u> | <u>Address</u> | <u>Sex</u> | <u>Salary</u> | Super_ssn | <u>Dno</u> |
|--------------|--------------|--------------|-----------|--------------|--------------------------|------------|---------------|-----------|------------|
| John | B | Smith | 123456789 | 1965-09-01 | 731 Fondren, Houston, TX | M | 30000 | 333445555 | 5 |
| Franklin | T | Wong | 333445555 | 1955-12-08 | 638 Voss, Houston, TX | M | 40000 | 888665555 | 5 |
| Ramesh | K | Narayan | 666884444 | 1962-09-15 | 975 Fire Oak, Humble, TX | M | 38000 | 333445555 | 5 |
| Joyce | A | English | 453453453 | 1972-07-31 | 5631 Rice, Houston, TX | F | 25000 | 333445555 | 5 |

Unspecified WHERE Clause and Use of *

- Specify an asterisk (*)
 - Query Q1D retrieves all the attributes of an EMPLOYEE and the attributes of the DEPARTMENT in which he or she works for every employee of the 'Research' department.
 - Query Q10A specifies the CROSS PRODUCT of the EMPLOYEE and DEPARTMENT tables.

Q1D: **SELECT** *

FROM EMPLOYEE, DEPARTMENT

WHERE Dname='Research' **AND** Dno=Dnumber;

Q10A: **SELECT** *

FROM EMPLOYEE, DEPARTMENT;

Tables as Sets in SQL

- SQL does not automatically eliminate duplicate tuples in query results
- For aggregate operations (will see sec 7.1.7) duplicates must be accounted for
- Use the keyword **DISTINCT** in the `SELECT` clause
 - Only distinct tuples –rows - remain in the result

Query 11. Retrieve the salary of every employee (Q11) and all distinct salary values (Q11A).

Q11: **SELECT** ALL Salary
 FROM EMPLOYEE;

Q11A: **SELECT** **DISTINCT** Salary
 FROM EMPLOYEE;

Tables as Sets in SQL

- Set operations
 - **UNION, EXCEPT** (difference), **INTERSECT**
 - Corresponding multiset operations: **UNION ALL, EXCEPT ALL, INTERSECT ALL**)
 - Type compatibility is needed for these operations to be valid

Query 4. Make a list of all project numbers for projects that involve an employee whose last name is 'Smith', either as a worker or as a manager of the department that controls the project.

```
Q4A: ( SELECT    DISTINCT Pnumber
      FROM      PROJECT, DEPARTMENT, EMPLOYEE
      WHERE     Dnum=Dnumber AND Mgr_ssn=Ssn
              AND Lname='Smith' )

      UNION

      ( SELECT    DISTINCT Pnumber
      FROM      PROJECT, WORKS_ON, EMPLOYEE
      WHERE     Pnumber=Pno AND Essn=Ssn
              AND Lname='Smith' );
```

Substring Pattern Matching and Arithmetic Operators

- **LIKE** comparison operator
 - Used for string **pattern matching**
 - % replaces an arbitrary number of zero or more characters
 - underscore (_) replaces a single character
 - Examples: **WHERE** Address **LIKE** '%Houston,TX%';
- Query 12. Retrieve all employees whose address is in Houston, Texas.

```
Q12:  SELECT Fname, Lname
        FROM EMPLOYEE
        WHERE Address LIKE '%Houston%';
```

| Results | | | Messages |
|---------|----------|---------|----------|
| | Fname | Lname | |
| 1 | John | Smith | |
| 2 | Franklin | Wong | |
| 3 | Joyce | English | |
| 4 | James | Borg | |
| 5 | Ahmad | Jabbar | |

Substring Pattern Matching and Arithmetic Operators

- **BETWEEN** comparison operator
- Ex: Q14 : WHERE (Salary **BETWEEN** 30000 **AND** 40000) **AND** Dno = 5;
- Query 14. Retrieve all employees in department 5 whose salary is between \$30,000 and \$40,000.
- **Q14**: SELECT * FROM EMPLOYEE WHERE
 (Salary BETWEEN 30000 AND 40000) AND Dno = 5;
- Alternate condition: The condition (Salary **BETWEEN** 30000 **AND** 40000) in Q14 is equivalent to the condition ((Salary >= 30000) **AND** (Salary <= 40000)).

Arithmetic Operations

- Standard arithmetic operators:
 - Addition (+), subtraction (−), multiplication (*), and division (/) may be included as a part of **SELECT**
- Query-13. Show the resulting salaries if every employee working on the 'ProductX' project is given a 10 percent raise.

**Q13: SELECT E.Fname, E.Lname, 1.1 * E.Salary AS Increased_sal
FROM EMPLOYEE AS E, WORKS_ON AS W, PROJECT AS P
WHERE E.Ssn=W.Essn AND W.Pno=P.Pnumber AND
P.Pname='ProductX';**

Ordering of Query Results

- Use **ORDER BY** clause
 - Keyword **DESC** to see result in a descending order of values
 - Keyword **ASC** to specify ascending order explicitly
 - Typically placed at the end of the query

```
ORDER BY D.Dname DESC, E.Lname ASC,  
E.Fname ASC
```

- **Q15**

```
SELECT D.Dname, E.Lname, E.Fname, P.Pname  
FROM DEPARTMENT AS D, EMPLOYEE AS E, WORKS_ON AS W, PROJECT AS P  
WHERE D.Dnumber = E.Dno AND E.Ssn = W.Essn AND W.Pno = P.Pnumber  
ORDER BY D.Dname, E.Lname, E.Fname;
```

Basic SQL Retrieval Query Block

- A simple retrieval query in SQL can consist of up to four clauses, but only first two—SELECT and FROM—are mandatory.
- The SELECT clause lists the attributes to be retrieved, and the FROM clause specifies all tables needed in the query. The WHERE clause identifies the conditions for selecting the rows from these tables, including join conditions if needed. ORDER BY specifies an order for displaying the results of a query.

```
SELECT      <attribute list>  
FROM        <table list>  
[ WHERE     <condition> ]  
[ ORDER BY <attribute list> ];
```

Additional Features of SQL

- Techniques for specifying complex retrieval queries.
(will see Ch.7)
- Writing programs in various programming languages that include SQL statements: Embedded and dynamic SQL, SQL/CLI (Call Level Interface) and its predecessor ODBC, SQL/PSM (Persistent Stored Module).
(will See Ch.10)
- Set of commands for specifying physical database design parameters, file structures for relations, and access paths, e.g., CREATE INDEX

Additional Features of SQL

- Transaction control commands.
- Specifying the granting and revoking of privileges to users (will see at Ch.30)
- Constructs for creating triggers (will see at Ch.10 and 26)
- Enhanced relational systems known as object-relational define relations as classes. Abstract data types (called User Defined Types- UDTs) are supported with CREATE TYPE
- New technologies such as XML (Ch.13) and OLAP (Ch.29) are added to versions of SQL

End Chapter Questions

- 6.1. How do the relations (tables) in SQL differ from the relations defined formally in Chapter 3? Discuss the other differences in terminology.
- 6.4. Describe the four clauses in the syntax of a simple SQL retrieval query. Show what type of constructs can be specified in each of the clauses. Which are required and which are optional?