

MTH522 Block: 1; Project: 2

Name: Pradyoth Singenahalli Prabhu

9. This exercise involves the Auto data set studied in the lab. Make sure that the missing values have been removed from the data.

Importing required libraries

```
In [1]: import pandas as pd
import string
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
plt.rcParams['figure.dpi'] = 100
alphabet = string.ascii_letters+string.punctuation
```

Importing data

```
In [2]: df = pd.read_csv("https://static1.squarespace.com/static/5ff2adbe3fe4fe33db9
```

Analysis

```
In [3]: df.columns
```

```
Out[3]: Index(['mpg', 'cylinders', 'displacement', 'horsepower', 'weight',
              'acceleration', 'year', 'origin', 'name'],
              dtype='object')
```

```
In [4]: df.head()
```

```
Out[4]:
```

	mpg	cylinders	displacement	horsepower	weight	acceleration	year	origin	name
0	18.0	8	307.0	130	3504	12.0	70	1	chevrolet chevelle malibu
1	15.0	8	350.0	165	3693	11.5	70	1	buick skylark 320
2	18.0	8	318.0	150	3436	11.0	70	1	plymouth satellite
3	16.0	8	304.0	150	3433	12.0	70	1	amc rebel sst
4	17.0	8	302.0	140	3449	10.5	70	1	ford torino

```
In [5]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 397 entries, 0 to 396
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   mpg              397 non-null   float64
1   cylinders         397 non-null   int64
2   displacement      397 non-null   float64
3   horsepower        397 non-null   object
4   weight            397 non-null   int64
5   acceleration      397 non-null   float64
6   year              397 non-null   int64
7   origin            397 non-null   int64
8   name              397 non-null   object
dtypes: float64(3), int64(4), object(2)
memory usage: 28.0+ KB
```

In [6]: `df.describe()`

Out[6]:

	mpg	cylinders	displacement	weight	acceleration	year	
count	397.000000	397.000000	397.000000	397.000000	397.000000	397.000000	397.0
mean	23.515869	5.458438	193.532746	2970.261965	15.555668	75.994962	1.5
std	7.825804	1.701577	104.379583	847.904119	2.749995	3.690005	0.8
min	9.000000	3.000000	68.000000	1613.000000	8.000000	70.000000	1.0
25%	17.500000	4.000000	104.000000	2223.000000	13.800000	73.000000	1.0
50%	23.000000	4.000000	146.000000	2800.000000	15.500000	76.000000	1.0
75%	29.000000	8.000000	262.000000	3609.000000	17.100000	79.000000	2.0
max	46.600000	8.000000	455.000000	5140.000000	24.800000	82.000000	3.0

Checking for missing values

In [7]: `df.isnull().values.any()`

Out[7]: False

In [8]: `df.isnull().sum()`

Out[8]:

```
mpg          0
cylinders     0
displacement  0
horsepower    0
weight        0
acceleration  0
year          0
origin        0
name          0
dtype: int64
```

Note: No missing values in the dataset.

In [9]: `df['horsepower'].unique()`

```
Out[9]: array(['130', '165', '150', '140', '198', '220', '215', '225', '190',
              '170', '160', '95', '97', '85', '88', '46', '87', '90', '113',
              '200', '210', '193', '?', '100', '105', '175', '153', '180', '110',
              '72', '86', '70', '76', '65', '69', '60', '80', '54', '208', '155',
              '112', '92', '145', '137', '158', '167', '94', '107', '230', '49',
              '75', '91', '122', '67', '83', '78', '52', '61', '93', '148',
              '129', '96', '71', '98', '115', '53', '81', '79', '120', '152',
              '102', '108', '68', '58', '149', '89', '63', '48', '66', '139',
              '103', '125', '133', '138', '135', '142', '77', '62', '132', '84',
              '64', '74', '116', '82'], dtype=object)
```

Removing symbol '?'

```
In [10]: df.horsepower.str.strip(alphabet).astype(bool).any()
```

```
Out[10]: True
```

```
In [11]: df = df[df.horsepower != '?']
```

```
In [12]: df.shape
```

```
Out[12]: (392, 9)
```

Converting data-type 'object' to 'int64'

```
In [13]: df["horsepower"] = pd.to_numeric(df["horsepower"])
df.dtypes
```

```
Out[13]: mpg                float64
cylinders                 int64
displacement             float64
horsepower               int64
weight                   int64
acceleration             float64
year                     int64
origin                   int64
name                     object
dtype: object
```

(a) Which of the predictors are quantitative, and which are qualitative?

Quantitative

1. mpg
2. cylinders
3. displacement
4. horsepower
5. weight
6. acceleration

Qualitative

1. year
2. origin

3. name

(b) What is the range of each quantitative predictor? You can answer this using the range() function.

dropping column 'year', 'origin' and 'name'

```
In [14]: df1 = df.copy()
df1 = df1.drop('name', axis=1)
df1 = df1.drop('year', axis=1)
df1 = df1.drop('origin', axis=1)
```

Range

```
In [15]: def get_range(df):
        for col in df.columns:
            print("Range of {0}: {1} to {2}".format(col, str(df[col].min()), str(df[col].max())))
```

```
In [16]: get_range(df1)

Range of mpg: 9.0 to 46.6
Range of cylinders: 3 to 8
Range of displacement: 68.0 to 455.0
Range of horsepower: 46 to 230
Range of weight: 1613 to 5140
Range of acceleration: 8.0 to 24.8
```

(c) What is the mean and standard deviation of each quantitative predictor?

Mean

```
In [17]: df1.mean(axis = 0)
```

```
Out[17]: mpg                23.445918
cylinders                5.471939
displacement            194.411990
horsepower              104.469388
weight                 2977.584184
acceleration             15.541327
dtype: float64
```

Standard Deviation

```
In [18]: df1.std()
```

```
Out[18]: mpg                7.805007
cylinders                1.705783
displacement            104.644004
horsepower              38.491160
weight                 849.402560
acceleration             2.758864
dtype: float64
```

(d) Now remove the 10th through 85th observations. What is the range, mean, and standard deviation of each predictor in the subset of the data that remains?

removing rows from 10 to 85

```
In [19]: df1.shape
```

```
Out[19]: (392, 6)
```

```
In [20]: df2 = df1.copy()
```

```
In [21]: df2.head()
```

```
Out[21]:
```

	mpg	cylinders	displacement	horsepower	weight	acceleration
0	18.0	8	307.0	130	3504	12.0
1	15.0	8	350.0	165	3693	11.5
2	18.0	8	318.0	150	3436	11.0
3	16.0	8	304.0	150	3433	12.0
4	17.0	8	302.0	140	3449	10.5

```
In [22]: df2.drop(axis=0, index=range(10,86), errors='ignore', inplace=True)
```

```
In [23]: df2.shape
```

```
Out[23]: (317, 6)
```

Mean

```
In [24]: df2.mean(axis = 0)
```

```
Out[24]:
```

mpg	24.374763
cylinders	5.381703
displacement	187.880126
horsepower	101.003155
weight	2938.854890
acceleration	15.704101
dtype:	float64

Standard Deviation

```
In [25]: df2.std()
```

```
Out[25]:
```

mpg	7.872565
cylinders	1.658135
displacement	100.169973
horsepower	36.003208
weight	811.640668
acceleration	2.719913
dtype:	float64

(e) Using the full data set, investigate the predictors


graphically, using scatterplots or other tools of your choice. Create some plots highlighting the relationships among the predictors. Comment on your findings.

```
In [26]: df['mpg'].value_counts()
```

```
Out[26]: 13.0    20
         14.0    19
         18.0    17
         15.0    16
         26.0    14
         ..
         31.9     1
         16.9     1
         18.2     1
         22.3     1
         44.0     1
         Name: mpg, Length: 127, dtype: int64
```

```
In [34]: sns.countplot(x = df['mpg'])
```

```
Out[34]: <AxesSubplot:xlabel='mpg', ylabel='count'>
```



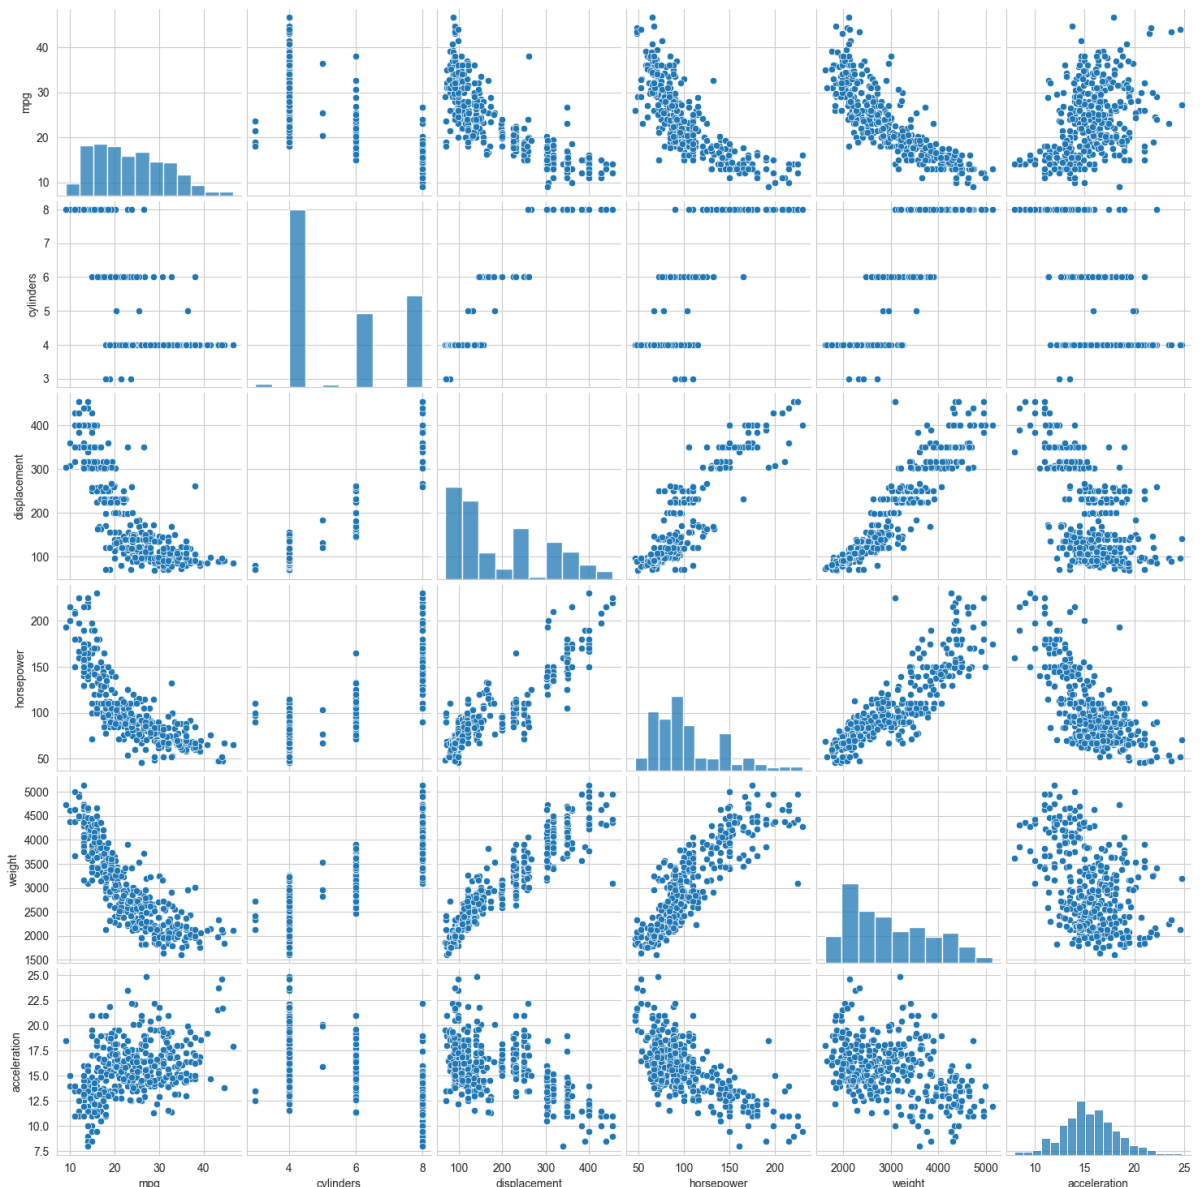
Observation:

Predicting 'mpg' is a regression problem.

Pair Plot

```
In [28]: plt.close()
```

```
sns.set_style("whitegrid")
sns.pairplot(df1)
plt.show()
```




Correlation Matrix

```
In [29]: plt.rcParams['figure.dpi'] = 150

corr_mat = df1.corr()
mask = np.array(corr_mat)
mask[np.tril_indices_from(mask)] = False
fig=plt.gcf()
fig.set_size_inches(4, 4)
sns.heatmap(data = corr_mat, mask = mask, square = True, annot = True, cbar
```

Out[29]: <AxesSubplot:>



The correlation matrix informs us on the relationship between quantitative predictors.


Observations:

1. Number of cylinders have positive correlation with all the quantitative predictors except 'acceleration'.
2. 'mpg' has positive relation with 'acceleration' and negative with all the other predictors.

Plotting with different features


```
In [30]: def box_plot(feature_x, target, dataframe = df1):
          sns.boxplot(x = target, y = feature_x, data = dataframe)
          plt.show()
```

```
In [31]: box_plot('mpg', 'year', df)
```


**Observation:**


1. We can see from the above bar plot, average 'mpg' increased from 1970 through 1980 and slightly decreased till 1982.
2. In the year 1980, 'mpg' increased drastically.
3. Bandwidth for year 1980 and 1982 is much higher compared to other years.

```
In [32]: box_plot('mpg', 'origin', df)
```

**Observation:**

1. 50 percentile of origin 3 is higher than 75 percentile of origin 2 and 25 percentile of origin 3 is more than 75 percentile of origin 1. So safely we can say 'mpg' in origin 3 automobiles is higher compared to others.

```
In [33]: box_plot('mpg', 'cylinders', df1)
```



Observations:

1. 'mpg' significantly increases when 4 cylinders are used compared to 3.
2. Lowest mpg is seen with 8 cylinders.
3. So higher the number of cylinders lower the 'mpg'. We can also witness in this in correctional matrix.

(f) Suppose that we wish to predict gas mileage (mpg) on the basis of the other variables. Do your plots suggest that any of the other variables might be useful in predicting mpg? Justify your answer.

1. From the pair plot and correlation matrix plotted above, we can see that 'mpg' has a negative correlation with cylinders, displacement, horsepower and weight. Also 'mpg' has positive correlation with acceleration.
2. Also from the above boxplots, we have observed that 'mpg' has a positive relation with year and origin.

So the above predictors can be used for predicting mpg.