

9. We will now consider the Boston housing data set, from the ISLR2 library.

```
In [1]: import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
```

```
In [2]: df_boston = pd.read_csv("/Volumes/work/sem_1/MTH522/data/Boston.csv")
df_boston.head()
```

```
Out[2]:
```

	Unnamed: 0	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	ptratio
0	1	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296	15.3
1	2	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242	17.8
2	3	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242	17.8
3	4	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222	18.7
4	5	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222	18.7

(a) Based on this data set, provide an estimate for the population mean of medv. Call this estimate $\hat{\mu}$.

```
In [3]: mu = df_boston['medv'].mean(axis=0)
print("Population Mean:", mu)
```

Population Mean: 22.532806324110677

(b) Provide an estimate of the standard error of $\hat{\mu}$. Interpret this result.

Standard error of the mean = s / \sqrt{n}

where:

s: sample standard deviation n: sample size

Source: <https://www.statology.org/standard-error-of-mean-python/>

```
In [4]: def std_err(data):  
        return np.std(data, ddof=1) / np.sqrt(np.size(data))  
  
print("Population Standard Error", std_err(df_boston['medv']))
```

Population Standard Error 0.40886114749753505

(c) Now estimate the standard error of $\hat{\mu}$ using the bootstrap. How does this compare to your answer from (b)?

Here below I have created 1000 samples of length 300.


```
In [5]: sample_means = list() # Empty list  
  
for _ in range(1000):  
    y = np.random.choice(df_boston['medv'], size=300, replace=True)  
    avg = np.mean(y)  
    sample_means.append(avg)
```

```
In [6]: print("Sample mean of mean:", np.mean(sample_means))  
        print("Standard deviation of sample means:", np.std(sample_means))
```

Sample mean of mean: 22.521797666666664
Standard deviation of sample means: 0.5413935552043243

```
In [7]: plt.hist(sample_means)  
        plt.ylabel("Count")  
        plt.xlabel("Mean")
```

Out[7]: Text(0.5, 0, 'Mean')



(d) Based on your bootstrap estimate from (c), provide a 95% confidence interval for the mean of medv. Compare it to the results obtained using `t.test(Boston$medv)`.

```
In [8]: std_err_boot = np.std(sample_means)
print(mu - 2 * std_err_boot, mu + 2 * std_err_boot)
21.450019213702028 23.615593434519326
```

(e) Based on this data set, provide an estimate, $\hat{\mu}_{med}$, for the median value of medv in the population.

```
In [9]: print("Population Median:", df_boston['medv'].median(axis=0))
Population Median: 21.2
```

(f) We now would like to estimate the standard error of $\hat{\mu}_{\text{med}}$. Unfortunately, there is no simple formula for computing the standard error of the median. Instead, estimate the standard error of the median using the bootstrap. Comment on your findings.

```
In [10]: sample_median = list() # Empty list


for _ in range(1000):
    y = np.random.choice(df_boston['medv'], size=300, replace=True)
    sample_median.append(np.median(y))
```

```
In [27]: print("Sample median of median:", np.median(sample_median))
print("Standard error of sample medians:", np.std(sample_median))
```

```
Sample median of median: 21.2
Standard error of sample medians: 0.46898954945712784
```

```
In [12]: plt.hist(sample_median)
plt.ylabel("Count")
plt.xlabel("Median")
```

```
Out[12]: Text(0.5, 0, 'Median')
```



(g) Based on this data set, provide an estimate for the tenth percentile of medv in Boston census tracts. Call this quantity $\hat{\mu}_{0.1}$. (You can use the `quantile()` function.)

```
In [13]: mu_10 = df_boston.medv.quantile(0.1)
print("Tenth percentile of medv is:", mu_10)
```

Tenth percentile of medv is: 12.75

(h) Use the bootstrap to estimate the standard error of $\hat{\mu}_{0.1}$. Comment on your findings.

```
In [24]: sample_quantile_10 = list()

for _ in range(1000):
    y = np.random.choice(df_boston['medv'], size=505, replace=True)
    sample_quantile_10.append(np.quantile(y, 0.1))
```

```
In [26]: print("Standard error of sample quantile of 10:", np.std(sample_quantile_
```

Standard error of sample quantile of 10: 0.5186870871729891

```
In [ ]:
```